

# HW3

Jadon Fowler jaf582 5778191

2022-10-18

## Instructions

Download a copy of this markdown. Change the `author:` tag above to have your name and NAU's ID. Fill in the file with your solution to the proposed problems. Knit your final document to PDF and submit it through BBLearn in the assignment [HW] Homework 3: review and cleaning by the end of the day on **Tuesday, October 18 (11:59 PM)**.

**Note:** to knit your homework to PDF, you need to have MikTeX installed. You can download and install MikTeX from [here](#). If you don't want to install MikTeX, you can knit to HTML, open the file in the browser and print the page to PDF.

For all the problems, please import the useful libraries in this chunk of code:

### Problem 1: Load the data

The dataset used in this assignment is the Automobile Data Set by Jeffrey C. Schlimmer. The goal is to use this data to predict **prices** of automobiles.

The following attributes are provided:

1. symboling: -3, -2, -1, 0, 1, 2, 3.
2. normalized-losses: continuous from 65 to 256.
3. make: alfa-romero, audi, bmw, chevrolet, dodge, honda, isuzu, jaguar, mazda, mercedes-benz, mercury,
4. fuel-type: diesel, gas.
5. aspiration: std, turbo.
6. num-of-doors: four, two.
7. body-style: hardtop, wagon, sedan, hatchback, convertible.
8. drive-wheels: 4wd, fwd, rwd.
9. engine-location: front, rear.
10. wheel-base: continuous from 86.6 to 120.9.
11. length: continuous from 141.1 to 208.1.
12. width: continuous from 60.3 to 72.3.
13. height: continuous from 47.8 to 59.8.
14. curb-weight: continuous from 1488 to 4066.
15. engine-type: dohc, dohcv, l, ohc, ohcf, ohcv, rotor.
16. num-of-cylinders: eight, five, four, six, three, twelve, two.
17. engine-size: continuous from 61 to 326.
18. fuel-system: 1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi.
19. bore: continuous from 2.54 to 3.94.
20. stroke: continuous from 2.07 to 4.17.
21. compression-ratio: continuous from 7 to 23.
22. horsepower: continuous from 48 to 288.

23. peak-rpm: continuous from 4150 to 6600.
24. city-mpg: continuous from 13 to 49.
25. highway-mpg: continuous from 16 to 54.
26. price: continuous from 5118 to 45400.

Load the .csv file and show the `head()`. What do you notice about variable names? Solve the issue and explain your solution.

The csv file doesn't have a line for column names. We can fix this by setting `col_names = FALSE` in `read_csv()` add we can add our own column names to the dataframe.

```
real.col.names <- c(
  "symboling",
  "normalized.losses",
  "make",
  "fuel.type",
  "aspiration",
  "num.of.doors",
  "body.style",
  "drive.wheels",
  "engine.location",
  "wheel.base",
  "length",
  "width",
  "height",
  "curb.weight",
  "engine.type",
  "num.of.cylinders",
  "engine.size",
  "fuel.system",
  "bore",
  "stroke",
  "compression.ratio",
  "horsepower",
  "peak.rpm",
  "city.mpg",
  "highway.mpg",
  "price"
)
auto <- as.data.frame(read_csv("auto.csv", col_names = FALSE))
```

```
## Rows: 205 Columns: 26
```

```
## -- Column specification -----
## Delimiter: ","
## chr (14): X2, X3, X4, X5, X6, X7, X8, X9, X15, X16, X18, X19, X22, X23
## dbl (12): X1, X10, X11, X12, X13, X14, X17, X20, X21, X24, X25, X26

##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
colnames(auto) <- real.col.names
head(auto)
```

```
##      symboling normalized.losses      make fuel.type aspiration num.of.doors
## 1      3      ? alfa-romero      gas      std      two
## 2      3      ? alfa-romero      gas      std      two
## 3      1      ? alfa-romero      gas      std      two
## 4      2      164      audi      gas      std      four
## 5      2      164      audi      gas      std      four
## 6      2      ?      audi      gas      std      two
##      body.style drive.wheels engine.location wheel.base length width height
## 1 convertible      rwd      front      88.6 168.8 64.1 48.8
## 2 convertible      rwd      front      88.6 168.8 64.1 48.8
## 3 hatchback      rwd      front      94.5 171.2 65.5 52.4
## 4      sedan      fwd      front      99.8 176.6 66.2 54.3
## 5      sedan      allwd      front      99.4 176.6 66.4 54.3
## 6      sedan      fwd      front      99.8 177.3 66.3 53.1
##      curb.weight engine.type num.of.cylinders engine.size fuel.system bore stroke
## 1      2548      dohc      four      130      mpfi 3.47 2.68
## 2      2548      dohc      four      130      mpfi 3.47 2.68
## 3      2823      ohcv      six      152      mpfi 2.68 3.47
## 4      2337      ohc      four      109      mpfi 3.19 3.40
## 5      2824      ohc      five      136      mpfi 3.19 3.40
## 6      2507      ohc      five      136      mpfi 3.19 3.40
##      compression.ratio horsepower peak.rpm city.mpg highway.mpg price
## 1      9.0      111      5000      21      27 13495
## 2      9.0      111      5000      21      27 16500
## 3      9.0      154      5000      19      26 16500
## 4      10.0      102      5500      24      30 13950
## 5      8.0      115      5500      18      22 17450
## 6      8.5      110      5500      19      25 15250
```

## Problem 2: Identifying the missing data

Inspect the dataset and the columns individually. Can you tell what codification is being used to address missing data?

Some columns use ?, some use -, and some use NA for missing data. We can normalize this by replacing weird values with NA.

```
auto$normalized.losses[auto$normalized.losses == '?'] <- NA
auto$num.of.doors[auto$num.of.doors == '-'] <- NA
auto$bore[auto$bore == '-'] <- NA
auto$horsepower[auto$horsepower == '?'] <- NA
auto$peak.rpm[auto$peak.rpm == '?'] <- NA
```

## Problem 3: Identifying typos

Several columns in this dataframe have typos (text that was misspelled by mistake). Identify and fix the typos. Explain your reasoning.

We can replace specific strings with other strings in the bad columns.


```
auto$make[auto$make == 'cherolet'] <- 'chevrolet'
auto$make[auto$make == 'plymouht'] <- 'plymouth'
```

#### Problem 4: dealing with the missing data

Plot the missing data, is there any columns/rows that we would benefit from dropping? Remember that the goal of the analysis is to predict the auto's prices. In case we have to deal with the missing values, which strategy would you use? Use functions and plots to demonstrate to argue on whether it would be interesting replace by the most frequent class or the mean value of the column. Explain your reasoning. Then, write a code that perform the changes you suggested.

Number of doors & number of cylinders can be replaced by the median Normalized.losses is missing over 40 items so it might be useful to just remove it. The rest of the columns can be replaced with the mean.

```
# plot na values vs column
na.values <- as.data.frame(sapply(auto, function (x) sum(is.na(x))))
na.values <- cbind(na.values, rownames(na.values))
rownames(na.values) <- NULL
colnames(na.values) <- c("na.count", "column")
ggplot(na.values) +
  geom_col(aes(x = na.count, y = column))
```



```

# replace NAs with medians
auto$num.of.cylinders[is.na(auto$num.of.cylinders)] <- median(auto$num.of.cylinders[!is.na(auto$num.of.cylinders)])
auto$num.of.doors[is.na(auto$num.of.doors)] <- median(auto$num.of.doors[!is.na(auto$num.of.doors)])

# replace NAs with means
auto$stroke[is.na(auto$stroke)] <- mean(as.numeric(auto$stroke[!is.na(auto$stroke)]))
auto$price[is.na(auto$price)] <- mean(as.numeric(auto$price[!is.na(auto$price)]))
auto$peak.rpm[is.na(auto$peak.rpm)] <- mean(as.numeric(auto$peak.rpm[!is.na(auto$peak.rpm)]))
auto$horsepower[is.na(auto$horsepower)] <- mean(as.numeric(auto$horsepower[!is.na(auto$horsepower)]))
auto$bore[is.na(auto$bore)] <- mean(as.numeric(auto$bore[!is.na(auto$bore)]))

```

### Problem 5: Type conversion

Look at the structure of your file. Are the variables represented in the appropriate data type? Convert the types for the ones that are not!

```

auto$symboling <- as.factor(auto$symboling)

number.lookup <- c(
  "one" = 1,
  "two" = 2,
  "three" = 3,
  "four" = 4,
  "five" = 5,
  "six" = 6,
  "seven" = 7,
  "eight" = 8,
  "nine" = 9,
  "ten" = 10
)

auto$num.of.cylinders <- number.lookup[auto$num.of.cylinders]
auto$num.of.doors <- number.lookup[auto$num.of.doors]

```

### Extra credit (10 points toward HW rubric): can you standardize the data?

Sometimes we get data from a source that uses a different measuring unit from what we expect. For example, in our in-class examples, we needed date in the mm/dd/yyyy format, but we got it in seconds. We had to transform the data into a more appropriate format/measuring unit. In data wrangling, this is called **Transformation**.

1. Transform the variables that stores the fuel consumption (variables 24. city-mpg and 25. highway-mpg) to store the values in liter per Km.

```

# convert mpg to kilometers per liter by dividing the value by 2.352
# and then turn that into liter per Km by flipping it
lpk <- function(x) 1 / (x / 2.352)
auto$city.mpg <- lpk(auto$city.mpg)
auto$highway.mpg <- lpk(auto$highway.mpg)

```

2. Transform the length measurements from inches to centimeters (variables 11. length, 12. width, 13. height, and 14. curb-weight).

```
auto$length <- auto$length * 2.54  
auto$width <- auto$width * 2.54  
auto$height <- auto$height * 2.54  
auto$curb.weight <- auto$curb.weight * 2.54
```